

A simple approximation to the average effect of the treatment on the treated in panel settings with selective enrollment

John Gardner*

August 2017

Abstract

This paper shows that, in panel settings when the probability of receiving a treatment is low, simple comparisons of outcomes between those who are treated in a given period and those who are untreated in that period, but receive the treatment in some other period, approximately identify the average effect of the treatment on the treated. Monte Carlo simulations reveal that this approximate selection-correction substantially reduces selection bias relative to naive comparisons of outcomes between treated and untreated units.

Keywords: Causal effect, treatment effect, selection bias.

JEL Codes: C1, C2, C13, C14.

1 Introduction

One of the central concerns of empirical economics is the possibility that self-selection on unobserved characteristics introduces bias into comparisons between treated and untreated individuals, invalidating the causal interpretation of those comparisons. This paper proposes a simple selection-correction that can be used to approximate the Average Effect of the Treatment on the Treated (ATT) in panel settings where the probability of receiving the treatment is low. The intuition for this approach is that, when the probability of receiving the treatment is low, individuals who are treated in any period are likely to differ systematically from those that never are. Consequently, comparing outcomes for those who are treated in a given period with those who are untreated in that period, but receive the treatment in some other period, approximates the average effect of the treatment among the treated population.

To be clear at the outset, this approach can only approximate the ATT, and can only do so well in panel settings where the probability of receiving the treatment is small (though it may be a useful indicator of the direction and severity of selection bias even at moderate treatment rates). However, in applications it is often the case that relatively few individuals receive the treatment in any given period. Moreover, the approximate selection-correction approach is simple and intuitive,

*366 Holman Hall, P.O. Box 1848, Department of Economics, University, MS 38677. Email: jrgardne@olemiss.edu.

and can be applied absent an instrument or other source of exogenous variation in treatment status, as well as in settings where the effect of unobserved characteristics on outcomes varies over time (in which traditional fixed-effects methods are invalid).

In the remainder of this paper, I provide theoretical motivation for the approximate selection-correction, illustrate its properties using a series of Monte Carlo simulations, demonstrate its application by using it to approximate the effect of Northward migration on Southern-born blacks' educational attainment, and conclude.

2 Approximating the ATT

Let Y_{1it} represent the outcome that individual $i \in \{1, \dots, N\}$ would experience at time $t \in \{1, \dots, T\}$ if they received a binary treatment ($D_{it} = 1$), and Y_{0it} the outcome they would experience absent the treatment (Rubin, 1974). Using this notation, i 's observed outcome can be expressed $Y_{it} = D_{it}Y_{1it} + (1 - D_{it})Y_{0it}$, and the time- t Average Effect of the Treatment on the Treated (ATT) can be expressed $ATT_t = E(Y_{1it} - Y_{0it} | D_{it} = 1)$. It is well known that when treatment status D_{it} is not randomly assigned, comparisons of observed outcomes between treated and untreated individuals identify a combination of the ATT and a bias term reflecting the possibility that treated individuals might have experienced systematically different outcomes even absent the treatment.

Suppose that potential outcomes (Y_{0it}, Y_{1it}) are independent of treatment status D_{it} conditional on an unobserved time-invariant characteristic U_i .¹ Suppose also that successive realizations of D_{it} are independent conditional on U_i , so that the U_i explain any serial correlation in treatment status. In this case, conditional on U_i , the bias term vanishes, and treated-untreated outcome comparisons identify U -specific causal effects:

$$\begin{aligned} E(Y_{it} | D_{it} = 1, U_i) - E(Y_{it} | D_{it} = 0, U_i) &= \underbrace{E(Y_{1t} - Y_{0t} | D_{it} = 1, U_i)}_{\text{ATT}} \\ &+ \underbrace{[E(Y_{0t} | D_{it} = 1, U_i) - E(Y_{0t} | D_{it} = 0, U_i)]}_{\text{Bias}} = E(Y_{1t} - Y_{0t} | D_{it} = 1, U_i). \end{aligned}$$

These strata-specific causal effects, in turn, can be aggregated according to the distribution of U_i among the treated population to identify the overall ATT. The argument that follows can be applied within covariate strata if potential outcomes are independent of treatment status conditional on U_i and observed covariates X_{it} .

This approach to identification is infeasible when U_i is unobserved. However, consider approximating the U_i by a binary variable $Q_i \in \{l, h\}$ with mass function $P(Q_i = h) = \pi$ and satisfying the restriction that $p_{lt} = P(D_{it} = 1 | Q_i = l) = 0$ for all t . In other words, approximate the unobserved heterogeneity by assuming that there are two types of individuals: those that will never receive the

¹The time-invariance assumption is for simplicity and can be relaxed.

treatment and those that might.²

Under this approximation, some individuals can be identified as members of group h with probability one. To see this, let $D_i = 1 \left(\sum_{t=1}^T D_{it} > 0 \right)$ be an indicator for whether individual i ever receives the treatment, and let $p_{ht} = P(D_{it} = 1 | Q_i = h)$ denote the probability that an individual belonging to group h receives the treatment at time t . The probability that an individual who ever receives the treatment belongs to group h is therefore simply

$$P(Q_i = h | D_i = 1) = \frac{\pi \left[1 - \prod_{t=1}^T (1 - p_{ht}) \right]}{\pi \left[1 - \prod_{t=1}^T (1 - p_{ht}) \right] + (1 - \pi) \left[1 - \prod_{t=1}^T (1 - p_{lt}) \right]} = 1.$$

Furthermore, the average effect of the treatment on the treated can be approximated as

$$\begin{aligned} ATT_t &= \int E(Y_{1it} - Y_{0it} | D_{it} = 1, U_i = u) dP(U_i = u | D_{it} = 1) \\ &= \int [E(Y_{it} | D_{it} = 1, U_i = u) - E(Y_{it} | D_{it} = 0, U_i = u)] dP(U_i = u | D_{it} = 1) \\ &\approx \sum_{q \in \{l, h\}} [E(Y_{it} | D_{it} = 1, Q_i = q) - E(Y_{it} | D_{it} = 0, Q_i = q)] P(Q_i = q | D_{it} = 1) \\ &= E(Y_{it} | D_{it} = 1, D_i = 1) - E(Y_{it} | D_{it} = 0, D_i = 1). \end{aligned}$$

In summary, under the two-type approximation described above, the time- t ATT is the same as the average effect of the treatment for group h , and can therefore be identified by comparing treated and untreated outcomes at time t , among those who receive the treatment in *any* period. Of course, the usefulness of this approach to identification depends on the quality of the approximation that it provides to the underlying ATT, which may be poor if the effects of, and likelihood of receiving, the treatment are very heterogeneous.³ However, the lower the probabilities $P(D_{it} = 1)$, $t \in \{1, \dots, T\}$, of receiving the treatment are, the better this approximation will be. This is because when there is meaningful unobserved heterogeneity in treatment probabilities but the overall population probability is low, it must be the case that a large proportion of the population receives the treatment with near-zero probability.⁴ Although this approximation argument is necessarily heuristic, I illustrate it with evidence from Monte Carlo simulations in the next section.

²In the terminology of Angrist, Imbens and Rubin (1996), this approximation separates the never-takers from the compliers, defiers, and always-takers.

³As I note in the appendix, this approximate selection-correction approach can also be viewed as an approximation to a formal latent-class analysis (see Leisch, 2004, for an introduction).

⁴E.g., consider a two-type model where a fraction π of the population enroll with probability p_h and $(1 - \pi)$ enroll with probability $p_l < p_h$. In broad terms, a small population probability $p = P(D_{it} = 1) = \pi p_h + (1 - \pi) p_l$ implies that either $p_l \approx 0$ and $p_h \approx p/\pi > p$ or $p_l \approx p_h \approx p$. In the latter case, unobserved heterogeneity would not introduce much selection bias, however, so only the former case is of interest here.

3 Monte Carlo simulations

I perform a number of Monte Carlo simulations to illustrate the properties of the approximation developed above. In each simulation, the data are generated by the following processes:

$$\begin{aligned} Y_{0it} &= bX_{it} + U_i + e_{0it} & Y_{1it} &= Y_{0it} + bX_{it} + U_i + e_{1it} \\ D_{it} &= 1(Y_{1it} - Y_{0it} > c) & Y_{it} &= D_{it}Y_{1it} + (1 - D_{it})Y_{0it}. \end{aligned}$$

I chose these processes to introduce selection bias into naive comparisons between treated and untreated individuals. In each simulation, the X_{it} are drawn from a standard normal distribution, and I vary the distributions from which the U_i , e_{0it} and e_{1it} are drawn, as well as the coefficient b on the observed covariate X_{it} , and the enrollment threshold c , as described below.

For each configuration, I simulate 1000 datasets consisting of observations on $(Y_{0it}, Y_{1it}, Y_{it}, D_{it}, X_{it})$ for $i \in \{1, \dots, 1000\}$ and $t \in \{1, \dots, 5\}$. For each simulate, I compute the sample average ATT as the sample analog of $E(Y_{1it} - Y_{0it} | D_{it} = 1)$, a naive estimate of the ATT as the coefficient on D_{it} from a regression of Y_{it} on D_{it} and X_{it} , and an approximately selection-corrected estimate of the ATT as the coefficient on D_{it} from the same regression estimated using the sample of individuals who ever receive the treatment. Note that the first of these objects can only be computed in a simulation setting where both treated and untreated outcomes are observed. Because the environment is stationary, I only report results for the final period. I summarize the simulation results graphically in Figure 1 and numerically in Table 1.⁵

In simulations (1)–(8), the (e_{0it}, e_{1it}) are drawn from independent standard normal distributions. For the first four of these simulations, I set the enrollment threshold c equal to 2.⁶ As the row labelled “Rate” in Table 1 shows, this generates population-average treatment probabilities of about 10% (per period). In simulation (1), I draw the U_i from a standard normal distribution and set the coefficient b on the observed covariate X_{it} to .1. The results are summarized in the first panel of Figure 1 and the first column of Table 1. A naive regression of observed outcomes on treatment status and the observed covariate overstates the average effect of the treatment on the treated considerably. In contrast, the approximately selection-corrected estimates are much closer to the true ATT, with the correction eliminating an average of about 73% $[=(4.07-3.03)/(4.07-2.64)]$ of the bias present in the naive estimate.

To illustrate the impact of *observed* heterogeneity on the quality of the approximation, in simulation (2) I set $b = 1$. As the row labelled “R sq.” in Table 1 shows, this increases the average R^2 from a regression of D_{it} on X_{it} from nearly zero to about .13. This additional source of variation in the probability of enrollment decreases the quality of the approximate selection-correction, which now eliminates an average of about 60% of the bias in the naive regression estimates. This pattern suggests that, when observed covariates matter, the approximate selection-correction should be

⁵In the plots in Figure 1, the bold lines represent the medians, the boxes represent the interquartile range, and the “whiskers” represent the minimum and maximum estimates (with outliers excluded according to R’s default algorithm).

⁶I chose this threshold because it was a whole number that generated a reasonably low treatment probability.

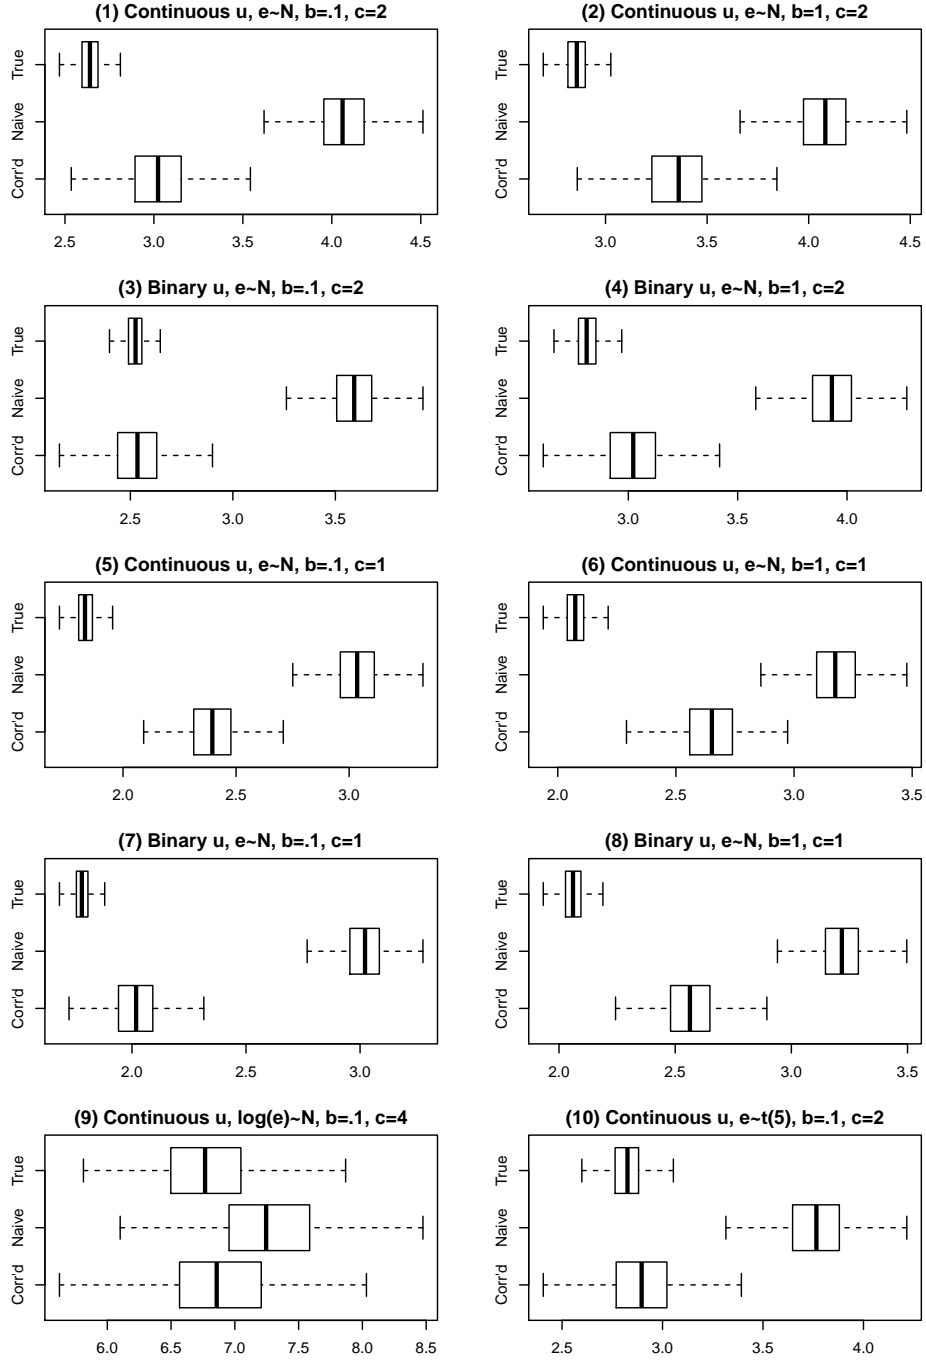


Figure 1: Distributions of ATTs from 1000 simulations. See main text for description.

applied within observed strata.

Table 1: Means and standard deviations from 1000 simulations

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)
True	2.64 (0.07)	2.86 (0.07)	2.53 (0.05)	2.81 (0.06)	1.83 (0.04)	2.08 (0.05)	1.78 (0.04)	2.06 (0.05)	6.79 (0.4)	2.83 (0.09)
Naive	4.07 (0.17)	4.08 (0.16)	3.59 (0.13)	3.93 (0.13)	3.04 (0.11)	3.18 (0.12)	3.02 (0.09)	3.22 (0.1)	7.28 (0.46)	3.77 (0.17)
Corrected	3.03 (0.19)	3.35 (0.18)	2.53 (0.14)	3.02 (0.15)	2.39 (0.12)	2.65 (0.13)	2.02 (0.11)	2.57 (0.12)	6.89 (0.47)	2.9 (0.19)
R sq.	0 (0)	0.13 (0.02)	0 (0)	0.14 (0.02)	0 (0)	0.19 (0.02)	0 (0)	0.18 (0.02)	0 (0)	0 (0)
Rate	0.08 (0.01)	0.12 (0.01)	0.08 (0.01)	0.13 (0.01)	0.24 (0.01)	0.28 (0.02)	0.26 (0.01)	0.29 (0.01)	0.09 (0.01)	0.1 (0.01)
Threshold	2	2	2	2	1	1	1	1	4	2
Het. type	Cts	Cts	Binary	Binary	Cts	Cts	Binary	Binary	Cts	Cts
b	0.1	1	0.1	1	0.1	1	0.1	1	0.1	0.1
Error distn.	N	N	N	N	N	N	N	N	LN(0,1)	t(5)

Notes—“Threshold” is the assumed value of c . “Het. type” indicates whether the U are drawn from a continuous (“Cts”) or equiprobable binary (“Binary”) distribution. “ b ” is the coefficient on the observable X . “Error distn.” indicates whether the (e_0, e_1) are drawn from normal (“N”), lognormal (“LN”), or t distributions with five degrees of freedom (“t(5)”).

For simulations (3) and (4), I repeat the analysis of simulations (1) and (2), drawing the $U_i \in \{-1, 1\}$ from an equiprobable binary distribution. The results show that the two-type selection correction provides a much better approximation to the true ATT when there is less variation in the underlying unobserved heterogeneity. Although the empiricist is unlikely to have prior information about such heterogeneity, these results illustrate the theoretical underpinnings of the approximation approach developed here.

To show how the quality of the approximate selection-correction varies with the probability of receiving the treatment, in simulations (5)-(8) I recreate simulations (1)-(4) with a lower enrollment threshold of $c = 1$, generating average treatment probabilities of about 25%. Consistent with the reasoning above, the performance of the approximate selection-correction declines as the probability of enrollment increases. In simulation (5), for example, the approximate correction only eliminates about half of the bias from the naive regression. Nevertheless, simulations (5)-(8) suggest that the approximate correction may be useful even at moderate treatment probabilities as an indicator of the direction and potential severity of bias due to selection on unobservables.

In simulations (9)–(10), I assess the sensitivity of the correction to the distribution from which the errors (e_{0it}, e_{1it}) are drawn. For brevity, these simulations use continuous U_i , set $b = .1$, and use enrollment thresholds that generate treatment rates close to 10% (unreported results from other configurations are similar). In simulation (9), these errors are drawn from independent lognormal distributions (i.e., the logs of the errors are standard normal). Though there is less absolute

selection bias in this setting (and more variation in the ATT), the selection correction removes over 90% of the bias that remains. In simulation (10), the errors are drawn from independent t distributions with five degrees of freedom. The fatter tails of this distribution increase the variance of the estimated true, naive, and corrected ATTs, which are otherwise similar to the normal case.

4 Application: The effect of the Great Migration on educational attainment

The African American Great Migration, during which millions of blacks left their birthplaces in the Southern US in favor of cities in the North, is thought to have had wide-reaching consequences for black-white relative economic progress in the 20th-century (see Tolnay, 2003, for a review). However, the possibility that the Great Migrants and their descendants were a select group that might have also experienced better socioeconomic outcomes in the South calls the casual interpretation of comparisons between migrants and non-migrants into question. To illustrate its application, I use the approximate selection-correction to estimate the effect of Northward migration on the educational attainment of migrants' children. My analysis draws on the Three-Generation Survey of Black American Families (Jackson and Tucker, 1997), an intergenerational dataset that reports socioeconomic outcomes and state of birth for three generations of a number of black families.⁷

The method developed in this paper is appropriate for this question for several reasons. First, to the extent that family-specific unobserved characteristics affect the educational attainment of parents and their children, they likely do so differently over time, violating the assumptions of fixed-effects methods. Second, the fractions of Southern-born individuals who migrated North are relatively small (e.g., only about 12% of second generation respondents whose parents were born in the South were themselves born in the North). Third, there are few clearly predetermined explanatory variables and no obviously exogenous sources of variation in migration with which to implement traditional regression or instrumental variables methods.

Table 2 presents the results. The first column reports estimates of a naive regression of educational attainment in years on indicators for being Northern-born and for being male among second-generation respondents whose parents were born in the South. The sample of second-generation respondents is used for comparability with the selection correction, which can only be implemented with that sample.⁸ The coefficient on Northern-born implies an estimated ATT of about 2 years. The second column reports approximately selection-corrected results from the same regression, estimated using a sample of second-generation respondents whose parents were born in the South but belonging to families that ever migrated to the North. Consistent with the notion

⁷I exclude observations for which the first-generation member of the family was born in a Northern state. Throughout my analysis, I use the Census Bureau's definition of the North and South. To account for the possibility that Southern-born members of the third generation (who are the youngest in the sample) might have migrated North after they were interviewed, I exclude third-generation respondents who were 18 or younger at the time of the survey.

⁸The data do not identify the birthplaces of first-generation respondents' parents, and the lack of return migration means that third-generation members of families who ever migrate North, but whose second-generation parents were Southern-born, must themselves be Northern-born.

that migrants and their children are positively selected on education, the implied ATT is lower, at about 1.16. For comparison, columns (3) and (4) report (family) fixed-effects estimates from pooled samples obtained using all three generations and only generations two and three. The implied ATTs of 3.58 and 1.29 exceed the approximate selection-corrected estimates, suggesting that a fixed-effects specification is inadequate in this intergenerational setting.

Table 2: The effect of parental migration on educational attainment

	Naive	Corr'd	FE (all)	FE (gen. 2 & 3)
Intercept	10.51*** (0.21)	11.56*** (0.41)	9.47*** (2.16)	11.35*** (2.02)
Northern-born	1.99*** (0.53)	1.16* (0.56)	3.58*** (0.45)	1.29** (0.50)
Male	-0.84* (0.38)	-1.45* (0.57)	0.40 (0.30)	0.30 (0.31)
R ²	0.05	0.10	0.36	0.53
Num. obs.	348	101	1044	696

Notes—Regressions of educational attainment in years on sex and an indicator for being born in the North among second-generation respondents with Southern-born parents. “Corr’d” denotes estimates obtained using the sample of families that ever migrate to the North. “FE” denotes a regression with family fixed effects (using either the entire sample or only members of the 2nd and 3rd generations). *** $p < 0.001$, ** $p < 0.01$, * $p < 0.05$.

5 Conclusion

This paper shows that, in panel settings when enrollment is selective, the ATT can be approximated using outcomes comparisons between treated and untreated units among the population that receives the treatment in any period. Monte Carlo simulations reveal that this approximate selection-correction can substantially reduce the bias present in naive comparisons of outcomes between treated and untreated individuals, particularly when the treatment probability is low, and may be useful as an indicator of the direction and severity of selection bias even when the treatment probability is modest. An application of this method suggests that, while Northward migration increased Southern-born blacks’ educational attainment, naive comparisons of education between the children of migrants and non-migrants overstate this effect.

References

Allman, Elizabeth S., Catherine Matias, and John A. Rhodes. 2009. “Identifiability of Parameters in Latent Structure Models with Many Observed Variables.” *The Annals of Statistics* 37 (6A): 3099-3132.

- Angrist, Joshua D., Guido W. Imbens, and Donald B. Rubin. 2007. "Identification of Causal Effects Using Instrumental Variables." *Journal of the American Statistical Association* 91 (434): 444-455.
- Bolck, Annabel, Marcel Croon, and Jacques Hagenars. 2004. "Estimating Latent Structure Models with Categorical Variables: One-Step versus Three-Step Estimators." *Political Analysis* 12 (1): 3-27.
- Jackson, James S., and M. Belinda Tucker. 1997. Three-Generation National Survey of Black American Families, 1979-1981. ICPSR09288-v2. Ann Arbor, MI: Inter-university Consortium for Political and Social Research [distributor].
- Tolnay, Stewart E. 2003. "The African American 'Great Migration' and Beyond." *Annual Review of Sociology* 29 (1): 209-32.
- Rubin, Donald. 1974. "Estimating Causal Effects of Treatments in Randomized and Nonrandomized Studies." *Journal of Educational Psychology* 66 (5): 688-701.
- Leisch, Friederich. 2004. "FlexMix: A General Framework for Finite Mixture Models and Latent Class Regression in R." *Journal of Statistical Software* 11 (1). <https://www.jstatsoft.org/article/view/v011i08>.

A Relationship to latent-class analysis

This approximate selection-correction approach can also be viewed as an approximation to a formal latent-class analysis, in which individuals are assumed to belong to one of a finite number of latent classes, the latent-class distribution and class-specific treatment probabilities are chosen to maximize the sample log-likelihood, the number of classes is chosen to maximize either the Bayesian or Akaike Information Criterion, and individuals are assigned to latent classes according to their probabilities of membership in each class, conditional on their observed characteristics (see Leisch, 2004, for an introduction). The advantages of the approximate approach include its conceptual and computational simplicity and the fact that it can be applied with only two time periods (while formal latent-class models are not identified with fewer than three periods, see Allman, Matias and Rhodes 2009).⁹

⁹Though latent-class analysis can be applied regardless of the treatment probabilities, when those probabilities are high, latent-class membership will be harder to identify from the sequences of treatment decisions, and treatment effect estimates obtained by comparing treated and untreated units assigned to the same classes will be ridden with greater classification error (this problem can be solved by modeling observed outcomes along with the treatment decisions, or by correcting for the resulting classification errors; see Bock, Croon and Hagenars 2004). At lower treatment rates, the approximate and formal latent-class methods produce similar estimates (as unreported simulation results show).